


```
In [4]: y = iris.target
y
```

```
Out[4]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
               0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
               0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
               1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
               1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
               2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
               2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
               2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2])
```

```
In [5]: df = pd.DataFrame(data=np.c_[X, y], columns=iris.feature_names + ['target'])
df
```

```
Out[5]:
```


	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
0	5.1	3.5	1.4	0.2	0.0
1	4.9	3.0	1.4	0.2	0.0
2	4.7	3.2	1.3	0.2	0.0
3	4.6	3.1	1.5	0.2	0.0
4	5.0	3.6	1.4	0.2	0.0
...
145	6.7	3.0	5.2	2.3	2.0
146	6.3	2.5	5.0	1.9	2.0
147	6.5	3.0	5.2	2.0	2.0
148	6.2	3.4	5.4	2.3	2.0
149	5.9	3.0	5.1	1.8	2.0

150 rows × 5 columns


```
In [6]: print(df.head())
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
0	5.1	3.5	1.4	0.2	0.0
1	4.9	3.0	1.4	0.2	0.0
2	4.7	3.2	1.3	0.2	0.0
3	4.6	3.1	1.5	0.2	0.0
4	5.0	3.6	1.4	0.2	0.0

```
In [7]: sns.pairplot(df, hue='target', markers=["o", "s", "D"])
plt.show()
```



```
sns.heatmap(df.corr(), annot=True)
plt.show()
```



```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
X_train, X_test, y_train, y_test
```

```
Out[9]: (array([[4.6, 3.6, 1. , 0.2],
                 [5.7, 4.4, 1.5, 0.4],
                 [6.7, 3.1, 4.4, 1.4],
                 [4.8, 3.4, 1.6, 0.2],
                 [4.4, 3.2, 1.3, 0.2],
                 [6.3, 2.5, 5. , 1.9],
                 [6.4, 3.2, 4.5, 1.5],
                 [5.2, 3.5, 1.5, 0.2],
                 [5. , 3.6, 1.4, 0.2],
                 [5.2, 4.1, 1.5, 0.1],
                 [5.8, 2.7, 5.1, 1.9],
                 [6. , 3.4, 4.5, 1.6],
                 [6.7, 3.1, 4.7, 1.5],
                 [5.4, 3.9, 1.3, 0.4],
                 [5.4, 3.7, 1.5, 0.2],
                 [5.5, 2.4, 3.7, 1. ],
                 [6.3, 2.8, 5.1, 1.5],
                 [6.4, 3.1, 5.5, 1.8],
                 [6.6, 3. , 4.4, 1.4],
                 [5.2, 3.6, 6.1, 0.5]]))
```

```
In [10]: scaler = StandardScaler()  
scaler
```

```
Out[10]: ▾ StandardScaler  
StandardScaler()
```

```
In [11]: X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)
```

```
In [12]: knn = KNeighborsClassifier(n_neighbors=5)  
knn.fit(X_train, y_train)
```

```
Out[12]: ▾ KNeighborsClassifier  
KNeighborsClassifier()
```

```
In [13]: y_pred = knn.predict(X_test)  
y_pred
```

```
Out[13]: array([1, 0, 2, 1, 1, 0, 1, 2, 1, 1, 2, 0, 0, 0, 0, 1, 2, 1, 1, 2, 0, 2,  
                0, 2, 2, 2, 2, 2, 0, 0])
```

```
In [14]: print("Confusion Matrix:")  
print(confusion_matrix(y_test, y_pred))
```

```
Confusion Matrix:  
[[10  0  0]  
 [ 0  9  0]  
 [ 0  0 11]]
```

```
In [15]: print("\nClassification Report:")  
print(classification_report(y_test, y_pred))
```


```
Classification Report:
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
In [16]: print("\nAccuracy Score:")  
print(accuracy_score(y_test, y_pred))
```

```
Accuracy Score:  
1.0
```

```
In [17]: plt.figure(figsize=(6,4))
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d', cmap='Blu
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
In [29]: from matplotlib.colors import ListedColormap

def plot_decision_boundaries(X, y, classifier, resolution=0.02):
    # Setup marker generator and color map
    markers = ('s', 'x', 'o', '^', 'v')
    colors = ('red', 'blue', 'lightgreen', 'gray', 'cyan')
    cmap = ListedColormap(colors[:len(np.unique(y))])

    # Plot the decision surface
    x1_min, x1_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    x2_min, x2_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, resolution),
                           np.arange(x2_min, x2_max, resolution))
    Z = classifier.predict(np.array([xx1.ravel(), xx2.ravel()]).T)
    Z = Z.reshape(xx1.shape)
    plt.contourf(xx1, xx2, Z, alpha=0.3, cmap=cmap)
    plt.xlim(xx1.min(), xx1.max())
    plt.ylim(xx2.min(), xx2.max())

    # Plot class samples
    for idx, cl in enumerate(np.unique(y)):
        plt.scatter(x=X[y == cl, 0],
                    y=X[y == cl, 1],
                    alpha=0.8,
                    c=colors[idx],
                    marker=markers[idx],
                    label=cl,
                    edgecolor='black')
```


```
In [30]: X_train_2D = X_train[:, :2]
X_test_2D = X_test[:, :2]
```

```
In [31]: knn_2D = KNeighborsClassifier(n_neighbors=5)
knn_2D.fit(X_train_2D, y_train)
```

```
Out[31]: ▾ KNeighborsClassifier
KNeighborsClassifier()
```


```
In [32]: plot_decision_boundaries(X_train_2D, y_train, classifier=knn_2D)
plt.xlabel('Sepal length (standardized)')
plt.ylabel('Sepal width (standardized)')
plt.legend(loc='upper left')
plt.title('KNN Decision Boundaries (Training set)')
plt.show()
```

C:\Users\Aditya Kudva\AppData\Local\Temp\ipykernel_22968\347347175.py:22:
 UserWarning: You passed a edgecolor/edgecolors ('black') for an unfilled m
 arker ('x'). Matplotlib is ignoring the edgecolor in favor of the facecol
 or. This behavior may change in the future.
 plt.scatter(x=X[y == c1, 0],




```
In [33]: plot_decision_boundaries(X_test_2D, y_test, classifier=knn_2D)
plt.xlabel('Sepal length (standardized)')
plt.ylabel('Sepal width (standardized)')
plt.legend(loc='upper left')
plt.title('KNN Decision Boundaries (Test set)')
plt.show()
```

C:\Users\Aditya Kudva\AppData\Local\Temp\ipykernel_22968\347347175.py:22: UserWarning: You passed a edgecolor/edgecolors ('black') for an unfilled marker ('x'). Matplotlib is ignoring the edgecolor in favor of the facecolor. This behavior may change in the future.
 plt.scatter(x=X[y == c1, 0],



```
In [34]: from sklearn.model_selection import cross_val_score

def evaluate_features(X, y):
    feature_names = iris.feature_names
    for i in range(X.shape[1]):
        X_new = np.delete(X, i, axis=1)
        scores = cross_val_score(KNeighborsClassifier(n_neighbors=5), X_new, y)
        print(f"Feature removed: {feature_names[i]}, Mean accuracy: {scores.mean()}")

evaluate_features(X, y)
```

Feature removed: sepal length (cm), Mean accuracy: 0.9667
 Feature removed: sepal width (cm), Mean accuracy: 0.9667
 Feature removed: petal length (cm), Mean accuracy: 0.9600
 Feature removed: petal width (cm), Mean accuracy: 0.9467

```
In [35]: from sklearn.model_selection import learning_curve


train_sizes, train_scores, test_scores = learning_curve(knn, X, y, cv=5,
                                                         train_sizes=np.linspace(0.1, 1.0, 10))

train_mean = np.mean(train_scores, axis=1)
train_std = np.std(train_scores, axis=1)
test_mean = np.mean(test_scores, axis=1)
test_std = np.std(test_scores, axis=1)

plt.plot(train_sizes, train_mean, label="Training score", color="r")
plt.plot(train_sizes, test_mean, label="Cross-validation score", color="g")

plt.fill_between(train_sizes, train_mean - train_std, train_mean + train_std, color="r")
plt.fill_between(train_sizes, test_mean - test_std, test_mean + test_std, color="g")

plt.title("Learning Curve")
plt.xlabel("Training Set Size")
plt.ylabel("Accuracy Score")
plt.legend(loc="best")
plt.show()
```



```
In [36]: from sklearn.model_selection import GridSearchCV

param_grid = {'n_neighbors': np.arange(1, 31)}
grid_search = GridSearchCV(KNeighborsClassifier(), param_grid, cv=5)
grid_search.fit(X_train, y_train)

print(f"Best parameters: {grid_search.best_params_}")
print(f"Best cross-validation score: {grid_search.best_score_:.4f}")
```

```
Best parameters: {'n_neighbors': 3}
Best cross-validation score: 0.9500
```

```
In [37]: from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize
from sklearn.multiclass import OneVsRestClassifier
from scipy import interp
from itertools import cycle
```

```
In [42]: y_test_bin = label_binarize(y_test, classes=[0, 1, 2])
n_classes = y_test_bin.shape[1]
```

```
In [43]: classifier = OneVsRestClassifier(KNeighborsClassifier(n_neighbors=5))
classifier
```

```
Out[43]:
```

```

> OneVsRestClassifier
> estimator: KNeighborsClassifier
    > KNeighborsClassifier

```


```
In [44]: y_score = classifier.fit(X_train, y_train).predict_proba(X_test)
y_score
```

```
Out[44]: array([[0. , 0.8, 0.2],
 [1. , 0. , 0. ],
 [0. , 0. , 1. ],
 [0. , 1. , 0. ],
 [0. , 1. , 0. ],
 [1. , 0. , 0. ],
 [0. , 1. , 0. ],
 [0. , 0. , 1. ],
 [0. , 0.6, 0.4],
 [0. , 1. , 0. ],
 [0. , 0. , 1. ],
 [1. , 0. , 0. ],
 [1. , 0. , 0. ],
 [1. , 0. , 0. ],
 [1. , 0. , 0. ],
 [0. , 0.8, 0.2],
 [0. , 0. , 1. ],
 [0. , 1. , 0. ],
 [0. , 1. , 0. ],
 [0. , 0. , 1. ],
 [1. , 0. , 0. ],
 [0. , 0.2, 0.8],
 [1. , 0. , 0. ],
 [0. , 0. , 1. ],
 [0. , 0. , 1. ],
 [0. , 0. , 1. ],
 [0. , 0.2, 0.8],
 [0. , 0. , 1. ],
 [1. , 0. , 0. ],
 [1. , 0. , 0. ]])
```

```
In [45]: fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_score[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])
```


```
In [46]: fpr["micro"], tpr["micro"], _ = roc_curve(y_test_bin.ravel(), y_score.ravel())
roc_auc["micro"] = auc(fpr["micro"], tpr["micro"])
```

```
In [47]: plt.figure()
colors = cycle(['aqua', 'darkorange', 'cornflowerblue'])
for i, color in zip(range(n_classes), colors):
    plt.plot(fpr[i], tpr[i], color=color, lw=2,
             label='ROC curve of class {0} (area = {1:0.2f})'
             ''.format(i, roc_auc[i]))
```



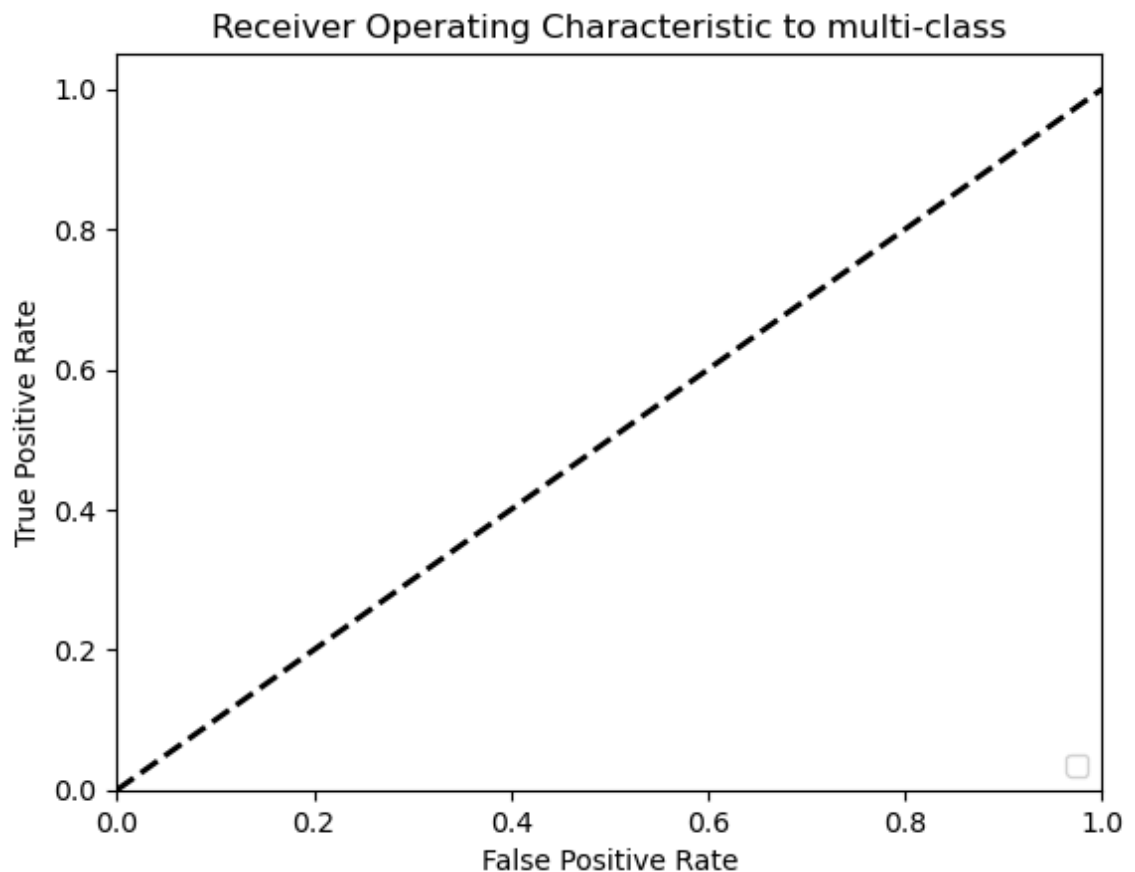
```
In [48]: plt.plot(fpr["micro"], tpr["micro"], color='deeppink', linestyle=':', linewidth=2,  
                label='micro-average ROC curve (area = {0:0.2f})'  
                ''.format(roc_auc["micro"]))
```

Out[48]: [



```
In [49]: plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic to multi-class')
plt.legend(loc="lower right")
plt.show()
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



In []: