

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, precision_score, recall_score,
```

```
In [4]: df = pd.read_csv('spam.csv', encoding='ISO-8859-1')
df
```

Out[4]:

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	NaN	NaN
...	...	...	...	...	...
5567	spam	This is the 2nd time we have tried 2 contact u...	NaN	NaN	NaN
5568	ham	Will Ì_b going to esplanade fr home?	NaN	NaN	NaN
5569	ham	Pity, * was in mood for that. So...any other s...	NaN	NaN	NaN
5570	ham	The guy did some bitching but I acted like i'd...	NaN	NaN	NaN
5571	ham	Rofl. Its true to its name	NaN	NaN	NaN

5572 rows × 5 columns

```
In [5]: print(df.head())
```

```

      v1                                v2 Unnamed: 2  \
0  ham  Go until jurong point, crazy.. Available only ...  NaN
1  ham                                Ok lar... Joking wif u oni...  NaN
2  spam  Free entry in 2 a wkly comp to win FA Cup fina...  NaN
3  ham  U dun say so early hor... U c already then say...  NaN
4  ham  Nah I don't think he goes to usf, he lives aro...  NaN

      Unnamed: 3  Unnamed: 4
0           NaN           NaN
1           NaN           NaN
2           NaN           NaN
3           NaN           NaN
4           NaN           NaN
```

```
In [7]: print(df.columns)
```

```
Index(['v1', 'v2', 'Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], dtype='object')
```

```
In [8]: df.rename(columns={'v1': 'label', 'v2': 'text'}, inplace=True)
```

```
In [9]: df['text'] = df['text'].str.lower().str.replace('[^\w\s]', '', regex=True)
```

```
In [10]: df['label'] = df['label'].map({'spam': 1, 'ham': 0})
```

```
In [11]: tfidf = TfidfVectorizer(stop_words='english', max_features=3000)
X = tfidf.fit_transform(df['text']).toarray()
y = df['label']
```

```
In [12]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
X_train, X_test, y_train, y_test
```

```
Out[12]: (array([[0., 0., 0., ..., 0., 0., 0.],
                 [0., 0., 0., ..., 0., 0., 0.],
                 [0., 0., 0., ..., 0., 0., 0.],
                 ...,
                 [0., 0., 0., ..., 0., 0., 0.],
                 [0., 0., 0., ..., 0., 0., 0.],
                 [0., 0., 0., ..., 0., 0., 0.]]) ,
          array([[0., 0., 0., ..., 0., 0., 0.],
                 [0., 0., 0., ..., 0., 0., 0.],
                 [0., 0., 0., ..., 0., 0., 0.],
                 ...,
                 [0., 0., 0., ..., 0., 0., 0.],
                 [0., 0., 0., ..., 0., 0., 0.],
                 [0., 0., 0., ..., 0., 0., 0.]]) ,
          1978      0
          3989      1
          3935      0
          4078      0
          4086      1
          ..
          3772      0
          5191      0
          5226      0
          5390      0
          860       0
          Name: label, Length: 4457, dtype: int64,
          3245      0
          944       0
          1044      1
          2484      0
          812       1
          ..
          4264      0
          2439      0
          5556      0
          4205      0
          4293      1
          Name: label, Length: 1115, dtype: int64)
```

```
In [13]: model = MultinomialNB()
model.fit(X_train, y_train)
```

```
Out[13]: ▾ MultinomialNB
MultinomialNB()
```

```
In [14]: y_pred = model.predict(X_test)
y_pred
```

```
Out[14]: array([0, 0, 1, ..., 0, 0, 1], dtype=int64)
```

```
In [15]: accuracy = accuracy_score(y_test, y_pred)
accuracy
```

```
Out[15]: 0.97847533632287
```

```
In [16]: precision = precision_score(y_test, y_pred)
precision
```

```
Out[16]: 1.0
```

```
In [17]: recall = recall_score(y_test, y_pred)
recall
```

```
Out[17]: 0.84
```


```
In [18]: f1 = f1_score(y_test, y_pred)
f1
```

```
Out[18]: 0.9130434782608696
```

```
In [19]: conf_matrix = confusion_matrix(y_test, y_pred)
conf_matrix
```


```
Out[19]: array([[965,  0],
               [ 24, 126]], dtype=int64)
```

```
In [20]: plt.figure(figsize=(10, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['Not Spam', 'Spam'], yticklabels=['Not Spam', 'Spam'], title='Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()
```



```
In [21]: fpr, tpr, _ = roc_curve(y_test, model.predict_proba(X_test)[: , 1])
roc_auc = auc(fpr, tpr)
```

```
In [22]: plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



```
In [23]: feature_names = tfidf.get_feature_names_out()
feature_log_probs = model.feature_log_prob_
```

```
In [24]: top_n = 20
top_spam_indices = np.argsort(feature_log_probs[1])[-top_n:]
top_ham_indices = np.argsort(feature_log_probs[0])[-top_n:]
```

```
In [25]: top_spam_features = [(feature_names[i], feature_log_probs[1][i]) for i in top_n]
top_spam_features
```

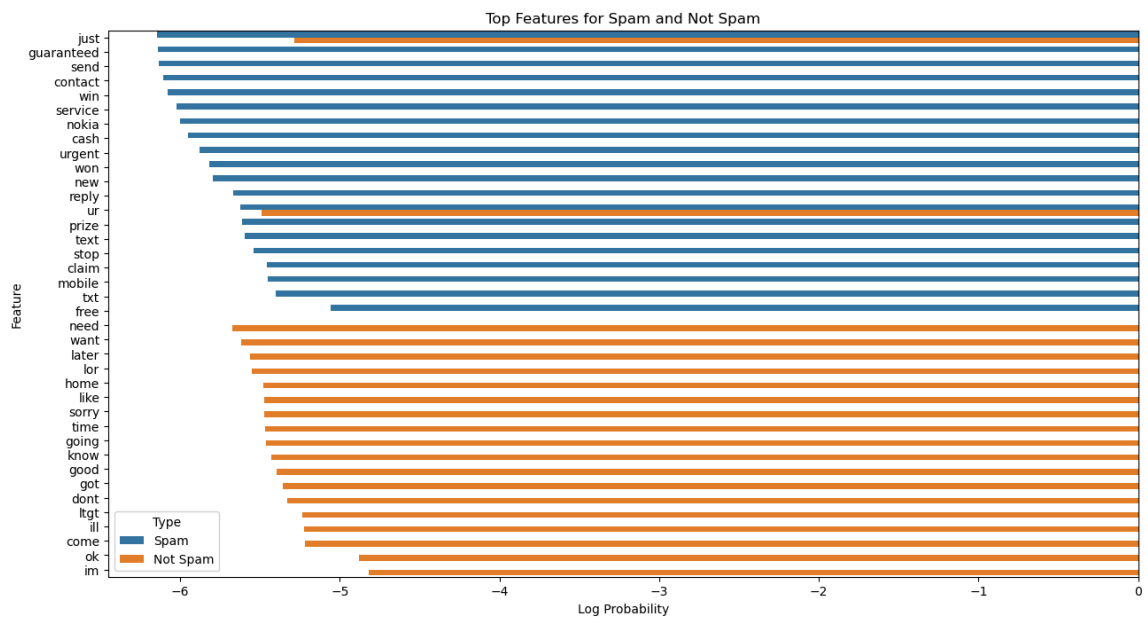
```
Out[25]: [('just', -6.143288062657874),
('guaranteed', -6.138908682040199),
('send', -6.136797507724539),
('contact', -6.104438562181709),
('win', -6.08025979848122),
('service', -6.022260868774975),
('nokia', -6.002258693589253),
('cash', -5.953783413457591),
('urgent', -5.878788966493407),
('won', -5.815846502468611),
('new', -5.798198987268488),
('reply', -5.668545003233329),
('ur', -5.6211056772013395),
('prize', -5.61307333115395),
('text', -5.595272962937164),
('stop', -5.541741976060994),
('claim', -5.4584015031284085),
('mobile', -5.449927956602329),
('txt', -5.401858933417513),
('free', -5.05998471623815)]
```

```
In [26]: top_ham_features = [(feature_names[i], feature_log_probs[0][i]) for i in top_n]
top_ham_features
```

```
Out[26]: [('need', -5.67553138973209),
('want', -5.62099126584992),
('later', -5.5619997755243755),
('lor', -5.549606264786631),
('ur', -5.492743062746311),
('home', -5.478612872669121),
('like', -5.4752903394950785),
('sorry', -5.474359407755339),
('time', -5.468271788022006),
('going', -5.464259018875763),
('know', -5.429825386578036),
('good', -5.397884672795875),
('got', -5.3566458745923065),
('dont', -5.33202939311373),
('just', -5.2847254016380605),
('ltgt', -5.233068600670436),
('ill', -5.2245966167844715),
('come', -5.219764470690236),
('ok', -4.880312054861962),
('im', -4.820912310519041)]
```

```
In [27]: top_features_df = pd.DataFrame(top_spam_features + top_ham_features, columns=feature_names)
top_features_df['Type'] = ['Spam'] * top_n + ['Not Spam'] * top_n
```


```
In [28]: plt.figure(figsize=(15, 8))
sns.barplot(x='Log Probability', y='Feature', hue='Type', data=top_features)
plt.title('Top Features for Spam and Not Spam')
plt.show()
```



```
In [29]: from sklearn.metrics import precision_recall_curve, average_precision_score
```

```
In [30]: y_proba = model.predict_proba(X_test)[: , 1]
precision, recall, _ = precision_recall_curve(y_test, y_proba)
average_precision = average_precision_score(y_test, y_proba)
```

```
In [31]: plt.figure(figsize=(10, 6))
plt.step(recall, precision, where='post', color='blue', alpha=0.7)
plt.fill_between(recall, precision, step='post', alpha=0.2, color='blue')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title(f'Precision-Recall Curve (AP = {average_precision:.2f})')
plt.show()
```




```
In [32]: conf_matrix_df = pd.DataFrame(conf_matrix, index=['True Not Spam', 'True Spam'], columns=['Predicted Not Spam', 'Predicted Spam'])
conf_matrix_df
```

Out[32]:

	Predicted Not Spam	Predicted Spam
True Not Spam	965	0
True Spam	24	126




```
In [33]: plt.figure(figsize=(10, 6))
sns.histplot(y_proba, kde=True, bins=30, color='blue')
plt.title('Distribution of Predicted Probabilities for Spam')
plt.xlabel('Predicted Probability')
plt.ylabel('Density')
plt.show()
```



```
In [34]: top_n = 20
top_spam_features = [(feature_names[i], feature_log_probs[1][i]) for i in top_n]
top_ham_features = [(feature_names[i], feature_log_probs[0][i]) for i in top_n]
```

```
In [35]: top_features_df = pd.DataFrame(top_spam_features + top_ham_features, columns=['Feature', 'Type'])
top_features_df['Type'] = ['Spam'] * top_n + ['Not Spam'] * top_n
```


```
In [36]: plt.figure(figsize=(15, 8))
sns.barplot(x='Log Probability', y='Feature', hue='Type', data=top_features)
plt.title('Top Features for Spam and Not Spam')
plt.show()
```



```
In [37]: plt.figure(figsize=(15, 8))
top_features_df = pd.DataFrame({
    'Feature': feature_names[top_spam_indices],
    'Spam Log Probability': feature_log_probs[1][top_spam_indices],
    'Not Spam Log Probability': feature_log_probs[0][top_spam_indices]
})
top_features_df = top_features_df.melt(id_vars='Feature', var_name='Type',
```

<Figure size 1500x800 with 0 Axes>

```
In [38]: sns.barplot(x='Log Probability', y='Feature', hue='Type', data=top_features_
plt.title('Feature Importance Comparison between Spam and Not Spam')
plt.show()
```



In [ ]: