

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [3]: data = pd.read_csv('sales.csv')
data
```

Out[3]:

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	
0	32298	CA-2012-124891	31/7/2012	31/7/2012	Same Day	RH-19495	Rick Hansen	Consumer	↑
1	26341	IN-2013-77878	5/2/2013	7/2/2013	Second Class	JR-16210	Justin Ritter	Corporate	Wc
2	25330	IN-2013-71249	17/10/2013	18/10/2013	First Class	CR-12730	Craig Reiter	Consumer	
3	13524	ES-2013-1579342	28/1/2013	30/1/2013	First Class	KM-16375	Katherine Murray	Home Office	
4	47221	SG-2013-4320	5/11/2013	6/11/2013	Same Day	RH-9495	Rick Hansen	Consumer	
...	
51285	29002	IN-2014-62366	19/6/2014	19/6/2014	Same Day	KE-16420	Katrina Edelman	Corporate	
51286	35398	US-2014-102288	20/6/2014	24/6/2014	Standard Class	ZC-21910	Zuschuss Carroll	Consumer	
51287	40470	US-2013-155768	2/12/2013	2/12/2013	Same Day	LB-16795	Laurel Beltran	Home Office	
51288	9596	MX-2012-140767	18/2/2012	22/2/2012	Standard Class	RB-19795	Ross Baird	Home Office	
51289	6147	MX-2012-134460	22/5/2012	26/5/2012	Second Class	MC-18100	Mick Crebagga	Consumer	

51290 rows × 24 columns



```
In [4]: data['Order Date'] = pd.to_datetime(data['Order Date'])
data['Ship Date'] = pd.to_datetime(data['Ship Date'])
```

C:\Users\Aditya Kudva\AppData\Local\Temp\ipykernel_31052\2887490760.py:1:
UserWarning: Parsing dates in DD/MM/YYYY format when dayfirst=False (the default) was specified. This may lead to inconsistently parsed dates! Specify a format to ensure consistent parsing.

data['Order Date'] = pd.to_datetime(data['Order Date'])
C:\Users\Aditya Kudva\AppData\Local\Temp\ipykernel_31052\2887490760.py:2:
UserWarning: Parsing dates in DD/MM/YYYY format when dayfirst=False (the default) was specified. This may lead to inconsistently parsed dates! Specify a format to ensure consistent parsing.

```
data['Ship Date'] = pd.to_datetime(data['Ship Date'])
```

```
In [5]: data['Order Year'] = data['Order Date'].dt.year
data['Order Month'] = data['Order Date'].dt.month
```

```
In [6]: features = ['Order Year', 'Order Month', 'Ship Mode', 'Segment', 'Market',
X = pd.get_dummies(data[features]) # Convert categorical variables into dummies
y = data['Sales']
```

```
In [7]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
X_train, X_test, y_train, y_test
```

```
Out[7]: (   Order Year  Order Month  Quantity  Discount  Shipping Cost  \
33162      2012           6         9      0.000         4.21
42206      2014          11         5      0.000         1.77
25603      2012           5         4      0.000         7.82
28126      2013           7         2      0.200         6.41
17208      2014           1         3      0.100        16.05
...         ...         ...         ...         ...         ...
11284      2013           9         2      0.500        28.64
44732      2012           4         2      0.500         1.28
38158      2014           4         3      0.000         2.69
860        2012           6         4      0.002        219.53
15795      2012           7         2      0.000        18.30
```

```
      Ship Mode_First Class  Ship Mode_Same Day  Ship Mode_Second Class
s \
33162                      0                  0
0
42206                      0                  0
0
25603                      1                  0
```


```
In [8]: model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[8]: LinearRegression
LinearRegression()
```


```
In [9]: y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f'Mean Squared Error: {mse:.2f}')
print(f'R^2 Score: {r2:.2f}')
```

Mean Squared Error: 62981.46
R^2 Score: 0.71

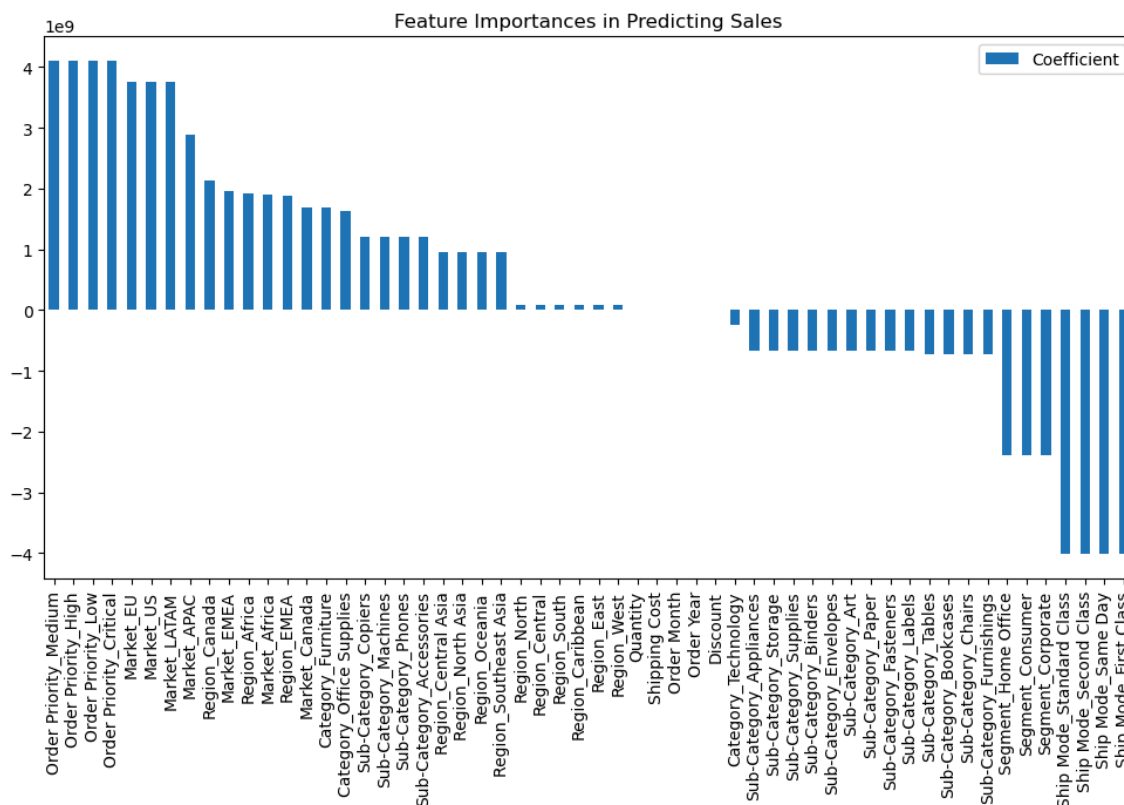
```
In [10]: plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, alpha=0.3)
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'k--', lw=4)
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Actual vs Predicted Sales')
plt.show()
```



```
In [11]: errors = y_test - y_pred
sns.histplot(errors, bins=25, kde=True)
plt.xlabel('Prediction Error')
plt.title('Prediction Error Distribution')
plt.show()
```



```
In [12]: feature_importance = pd.DataFrame(model.coef_, index=X.columns, columns=['Coefficient'])
feature_importance.sort_values(by='Coefficient', ascending=False, inplace=True)
feature_importance.plot(kind='bar', figsize=(12, 6))
plt.title('Feature Importances in Predicting Sales')
plt.show()
```



```
In [13]: sns.pairplot(data)
plt.suptitle('Pair Plot of Sales and Selected Features', verticalalignment=
plt.show())
```

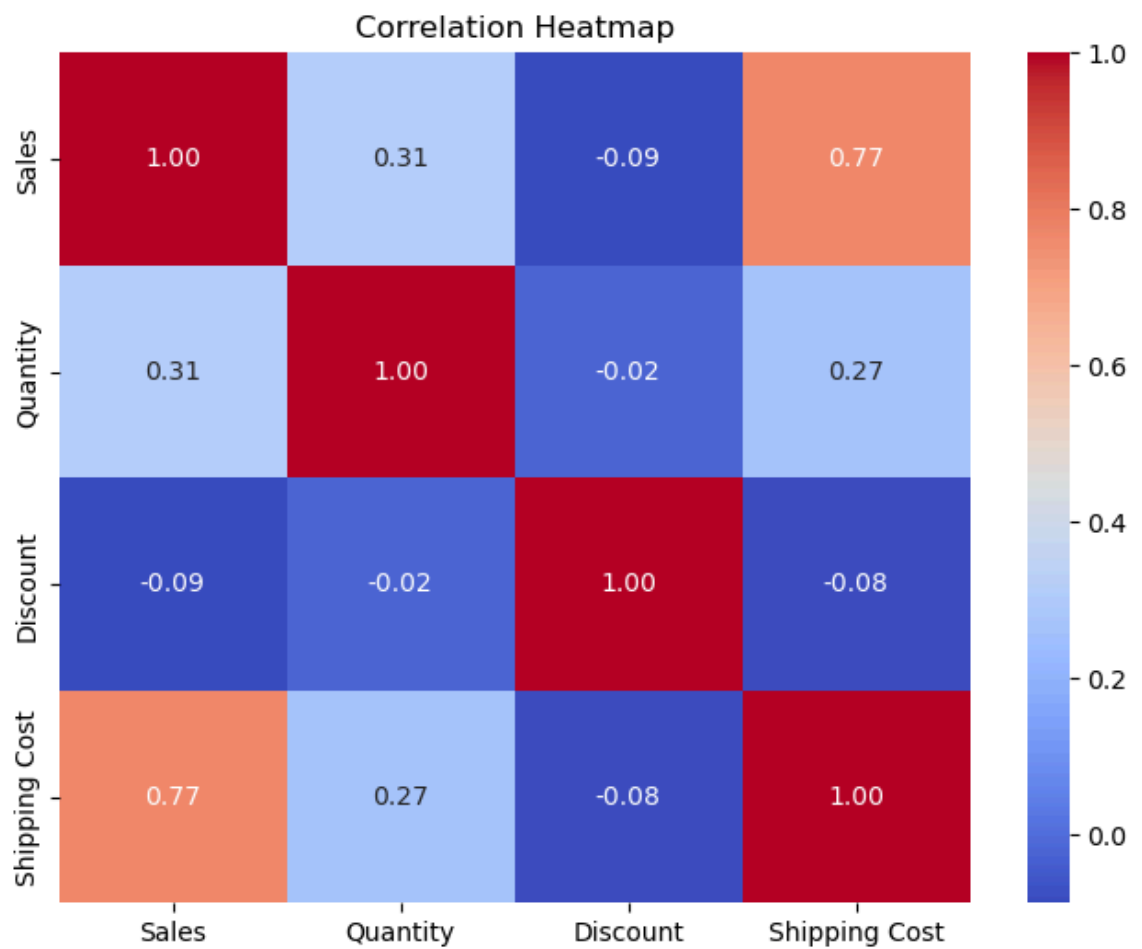


```
In [14]: correlation_matrix = data[['Sales', 'Quantity', 'Discount', 'Shipping Cost']]
correlation_matrix
```

Out[14]:

	Sales	Quantity	Discount	Shipping Cost
Sales	1.000000	0.313577	-0.086722	0.768073
Quantity	0.313577	1.000000	-0.019875	0.272649
Discount	-0.086722	-0.019875	1.000000	-0.079056
Shipping Cost	0.768073	0.272649	-0.079056	1.000000

```
In [15]: plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, fmt=".2f", cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```




```
In [16]: import plotly.express as px
```



```
In [17]: fig = px.line(data, x='Order Date', y='Sales', title='Sales Over Time', labels={'x': 'Order Date', 'y': 'Sales'})
fig.update_xaxes(rangeslider_visible=True)
fig.show()
```

Sales Over Time




```
In [18]: from sklearn.ensemble import RandomForestRegressor
```

```
In [19]: model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train, y_train)
```

```
Out[19]: ▼      RandomForestRegressor
          RandomForestRegressor(random_state=42)
```

```
In [20]: importances = model.feature_importances_
indices = np.argsort(importances)[::-1]
```


```
In [21]: plt.figure(figsize=(12, 6))
plt.title('Feature Importances')
plt.bar(range(X_train.shape[1]), importances[indices], align='center')
plt.xticks(range(X_train.shape[1]), X_train.columns[indices], rotation=90)
plt.xlim([-1, X_train.shape[1]])
plt.show()
```



```
In [22]: residuals = y_test - y_pred
residuals
```

```
Out[22]: 49728    -22.035868
45547     92.248218
15664     43.484829
40561    -34.142389
49426     43.925078
...
46156     55.910974
18754    -43.388948
29983    -32.241150
24864    -84.280607
12001     73.362399
Name: Sales, Length: 12823, dtype: float64
```

```
In [23]: plt.figure(figsize=(10, 6))
sns.residplot(x=y_pred, y=residuals, lowess=True, line_kws={'color': 'red'},
plt.title('Residuals vs. Predicted')
plt.xlabel('Predicted Sales')
plt.ylabel('Residuals')
plt.show()
```



In []: